

Gradient Descent


- What is a Gradient in math
- What is Gradient Descent
- Gradient Descent Algorithm
- Why Gradient Descent

What is Gradient in math


A gradient is a **vector** of partial derivatives for a **multivariate** function $f(x_1, x_2, \dots, x_n)$ with respect for each variable x_i

$$\nabla f = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right)$$


- The Gradient vector points in the direction of the **maximum** change (increase), formally **steepest ascent**
- Its magnitude is equal to the **maximum rate** of change



- A zero gradient $\nabla f = 0$ indicates a **critical point**



- The Gradient measures the "slop" in all directions our point here is on a *flat surface* the gradient vector is zero



- Unlike here where there is a slight slop the gradient vector ∇f points to the direction of the **steepest ascent**

Gradient Rules

- Product Rule

$$\nabla(f \cdot g) = f \nabla g + g \nabla f$$

- Quotient Rule

$$\nabla \left(\frac{f}{g} \right) = \frac{g \nabla f - f \nabla g}{g^2}$$

- Gradient of a Norm

$$\nabla \|x\| = \frac{x}{\|x\|}$$


- Directional Derivative Connection

$$D_{\hat{u}} G = \nabla G \cdot \hat{u}$$

What is Gradient Descent ?

Simply its an **optimization** algorithm used to **train machine learning** models, by optimizing the **parameters** of the model **iteratively** to find the **global minimum** of a loss function curve

More formally the Gradient Descent is an algorithm that essentially computes the gradient ∇ of the cost loss function (MSE in Linear Regression) $*h*$



- That is the lowest point in a function curve

Intuition :

Imagine a person(**gradient descent algorithm**) is stuck in a foggy mountain(**loss function curve**) and he is trying to get down (finding the **global minimum**). Therefore the person need to use local information and **calculations** and what's visible to descent down, using the gradient descent which says look at your current position and goes into the direction of the steepest descent

- **Fog** → Limited, local information
- **Slop** → Gradient
- **Step size** → Learning rate

Gradient Descent Algorithm

Generally the Gradient Descent follows These steps :

1. Initialize θ (randomly)
2. While not converged :
 1. Computer gradient : $\nabla J(\theta)$
 2. Update parameters : $\theta^+ = \theta^- - \alpha \nabla_{\theta} J(\theta)$
 3. Check convergence **optional**
3. Return optimized θ

Note :

- $J(\theta)$ is a the **Loss Function**
- α is the **Learning Rate** which is the step size
- **Batch size** → its the trade off between speed and stability

Iterations vs Epochs

An **Epoch** → is going through all the batches

- In **SGD** the Epoch = 1, cause we have one batch of the training example

An **iteration** is the number of **iterations** to complete one **Epoch**

- In **SGD** for example we need 5 iterations to complete one **Epoch** which is one batch of training example

Types Of Gradient Descent

1. Batch Gradient Descent
2. Stochastic Gradient Descent **SGD**
3. Mini-Batch Gradient Descent

Batch Gradient Descent

This is the **Vanilla Gradient Descent**, and works by :

- Using all the data set to compute the gradient of the cost function $\nabla J(\theta)$ in each iteration
- In simple terms it calculate all the gradients for every training example and take the **mean gradient** to make just one step for one iteration

Pros :

Its a *Slow but perfect approach* and **deterministic** which means we will get the same values each time


Down sides :

- Its **Computationally** expensive for large-medium data sets
- Can get stuck in **local minima** or **saddle points**

$$\theta^+ = \theta^- - \alpha \nabla_{\theta} J(\theta)$$

Stochastic Gradient Descent

Consider this the *Lazy Gradient Descent* which consider a random training examples from the data sets per iteration to compute the gradient



- Take this loss function for example using the **Full Batch Gradient Descent** it might take it a very long time to escape the local minima thinking its the minimum value for the loss function(cost function)

Pros :

- Here where the **Stochastic Gradient** can be useful cause its fluctuate due to randomness it escapes the **local minima** and results in faster updates which can be very effective in large data sets

Down sides :

- May never reach exactly the minimum of the cost function

- Also suffers from high variance

$$\theta^+ = \theta^- - \eta \nabla_{\theta} J(\theta; x^i, y^i)$$

Where (x^i, y^i) is a single training example

Mini-Batch Gradient Descent

Its a mixed approach between *Full Gradient Descent* and *Stochastic Gradient Descent* where instead of taking one example (x^i, y^i) from the training data, we take a small batch typically around 32, 64, 128, 256 . . .

Pros :

- Its a perfect mix between stability and efficiency, **Works faster than Full batch GD and less variance than SGD**
- Works well with GPU parallelization

Down sides :

- Just requires fine tuning the batch size

$$\theta^+ = \theta^- - \eta \nabla_{\theta} J(\theta; B)$$

Where B is the mini-batch of the training examples